

# Contents

<b>Kubernetes Learning Guide</b>	<b>3</b>
Who Is This For . . . . .	3
□ Kubernetes Philosophy . . . . .	3
Kubernetes Architecture (Simplified) . . . . .	5
Kubernetes Resources . . . . .	5
K3s Setup . . . . .	6
Learning Path . . . . .	7
Success Tips . . . . .	7
Let's Begin! . . . . .	8
Learning Objectives . . . . .	8
What is a Pod? . . . . .	8
Big Picture: Where Does Pod Fit? . . . . .	9
Hands-on Exercises . . . . .	9
Exam Practice . . . . .	13
Common Mistakes . . . . .	13
□ Cleanup . . . . .	13
What We Learned . . . . .	14
Next Steps . . . . .	14
Learning Objectives . . . . .	14
Multi-Container Patterns . . . . .	14
Hands-on Exercises . . . . .	14
Exam Practice . . . . .	19
□ Cleanup . . . . .	20
What We Learned . . . . .	20
Learning Objectives . . . . .	20
What is a Deployment? . . . . .	20
Hands-on Exercises . . . . .	20
Exam Practice . . . . .	25
□ Cleanup . . . . .	26
What We Learned . . . . .	26
Learning Objectives . . . . .	26
What is a Service? . . . . .	26
Hands-on Exercises . . . . .	27
Exam Practice . . . . .	30
□ Cleanup . . . . .	31
What We Learned . . . . .	31
Learning Objectives . . . . .	31
ConfigMap vs Secret . . . . .	31
ConfigMap Exercises . . . . .	31
Secret Exercises . . . . .	33
Exam Practice . . . . .	35
□ Cleanup . . . . .	35
What We Learned . . . . .	35
Learning Objectives . . . . .	35
Volume Types . . . . .	36
Hands-on Exercises . . . . .	36
Exam Practice . . . . .	39
□ Cleanup . . . . .	39
What We Learned . . . . .	40
Learning Objectives . . . . .	40
Job vs CronJob . . . . .	40
Job Exercises . . . . .	40

CronJob Exercises . . . . .	42
Exam Practice . . . . .	43
□ Cleanup . . . . .	43
What We Learned . . . . .	43
Learning Objectives . . . . .	44
Probe Types . . . . .	44
Hands-on Exercises . . . . .	45
Exam Practice . . . . .	47
□ Cleanup . . . . .	48
What We Learned . . . . .	48
Learning Objectives . . . . .	48
Requests vs Limits . . . . .	48
Hands-on Exercises . . . . .	49
Exam Practice . . . . .	51
□ Cleanup . . . . .	52
What We Learned . . . . .	52
Learning Objectives . . . . .	52
What is NetworkPolicy? . . . . .	52
Hands-on Exercises . . . . .	53
Exam Practice . . . . .	55
□ Cleanup . . . . .	56
What We Learned . . . . .	56
Learning Objectives . . . . .	56
What is SecurityContext? . . . . .	56
SecurityContext Exercises . . . . .	56
ServiceAccount Exercises . . . . .	58
Exam Practice . . . . .	59
□ Cleanup . . . . .	60
What We Learned . . . . .	60
Learning Objectives . . . . .	60
What is Ingress? . . . . .	60
Hands-on Exercises . . . . .	61
Exam Practice . . . . .	64
□ Cleanup . . . . .	64
What We Learned . . . . .	64
Learning Objectives . . . . .	64
Debugging Flow . . . . .	65
Hands-on Exercises . . . . .	65
Common Issues Cheatsheet . . . . .	67
Exam Practice . . . . .	67
□ Cleanup . . . . .	68
What We Learned . . . . .	68
Learning Objectives . . . . .	68
What is Helm? . . . . .	68
Hands-on Exercises . . . . .	69
Helm Commands Summary . . . . .	70
Exam Practice . . . . .	71
□ Cleanup . . . . .	71
What We Learned . . . . .	71
Learning Objectives . . . . .	72
Deployment Strategies . . . . .	72
Blue-Green Deployment . . . . .	72
Canary Deployment . . . . .	74
Exam Practice . . . . .	75

□ Cleanup . . . . .	75
What We Learned . . . . .	75
Learning Objectives . . . . .	75
What is Kustomize? . . . . .	76
Hands-on Exercises . . . . .	77
Kustomization.yaml Options . . . . .	79
Exam Practice . . . . .	80
□ Cleanup . . . . .	80
What We Learned . . . . .	80
Learning Objectives . . . . .	80
Container Structure . . . . .	80
Hands-on Exercises . . . . .	80
Exam Practice . . . . .	84
□ Cleanup . . . . .	85
What We Learned . . . . .	85

# Kubernetes Learning Guide

## Who Is This For?

This repo is for anyone who wants to **learn Kubernetes from scratch** and **prepare for the CKAD exam**.

## Prerequisites

- Basic Linux terminal knowledge (`cd`, `ls`, `cat`, `vim`)
- Know what Docker is (don't need to have used it)
- A running K3s cluster (setup instructions below)

## Target Audience

- Kubernetes beginners
  - CKAD exam candidates
  - DevOps/Cloud career seekers
- 

## □ Kubernetes Philosophy

### The Problem: Who Will Manage the Containers?


Docker is great, but... - How do you manage 100 containers? - If one crashes, does it auto-restart? - How does it scale when traffic increases? - How do containers find each other?

### The Solution: Kubernetes (Container Orchestrator)

**Kubernetes Core Principle:** > “Desired State” → Kubernetes makes it happen and maintains it

You: “I want 3 nginx pods” Kubernetes: “Done. If one dies, I’ll create a new one.”

---



- FAQ
- Who manages?
- What if one dies?
- How to scale?
- Networking?

Figure 1: Mermaid Diagram




Figure 2: Mermaid Diagram

## Kubernetes Architecture (Simplified)




Figure 3: Mermaid Diagram

### Simple Explanation

Component	Role	Real-World Analogy
<b>API Server</b>	Receives all requests	Company reception
<b>etcd</b>	Stores all data	Company database
<b>Scheduler</b>	Places pods on nodes	HR department
<b>Controller Manager</b>	Ensures everything works	Manager
<b>Kubelet</b>	Runs pods on node	Employee
<b>Kube Proxy</b>	Manages network traffic	IT department

## Kubernetes Resources

Everything in Kubernetes is a “resource”. Here’s what you’ll learn:

### Resource Hierarchy

```
Deployment
  ReplicaSet
    Pod
      Container(s)
        Image
```




Figure 4: Mermaid Diagram

## K3s Setup

K3s is a lightweight Kubernetes distribution. Perfect for learning!

### Linux Installation

```

# Install K3s (one command!)
curl -sfL https://get.k3s.io | sh -

# Check status
sudo systemctl status k3s

# Configure kubectl
mkdir -p ~/.kube
sudo cp /etc/rancher/k3s/k3s.yaml ~/.kube/config
sudo chown $(id -u):$(id -g) ~/.kube/config

# Test
kubectl get nodes

```

### WSL2 Installation

```

# WSL2 needs systemd
# Edit /etc/wsl.conf:
sudo nano /etc/wsl.conf
# Add:
# [boot]
# systemd=true

# Restart WSL (in PowerShell):
# wsl --shutdown

```

```
# Then install K3s
curl -sfL https://get.k3s.io | sh -
```

---

## Learning Path




Figure 5: Mermaid Diagram

### Week 1: Learn the Basics

1. **Lab 01** - What is a Pod? How to create one?
2. **Lab 02** - Multiple containers
3. **Lab 03** - Deployment for pod management
4. **Lab 04** - Service for access

### Week 2: Configuration

5. **Lab 05** - ConfigMap and Secret
6. **Lab 06** - Persistent storage (Volume)
7. **Lab 09** - CPU/Memory limits

### Week 3: Advanced Topics

8. **Lab 07** - Job and CronJob
9. **Lab 08** - Health checks (Probes)
10. **Lab 10** - Network security
11. **Lab 11** - Security settings

### Week 4: Tools and Deployment

12. **Lab 12** - Ingress (HTTP traffic)
  13. **Lab 13** - Debugging
  14. **Lab 14** - Helm (package management)
  15. **Lab 15** - Blue-Green/Canary
  16. **Lab 16** - Kustomize
  17. **Lab 17** - Dockerfile
- 

## Success Tips

### To Pass the Exam

1. **Practice, Practice, Practice** - Theory isn't enough
2. **Master kubectl** - Everything is CLI

3. **Don't memorize YAML** - Use `--dry-run=client -o yaml`
4. **Time management** - 2 hours, 15-20 questions
5. **Use kubernetes.io** - You have access during the exam!

## Learn the Aliases

```
alias k=kubectl
export do="--dry-run=client -o yaml"
export now="--force --grace-period=0"
```

---

## Let's Begin!

### Lab 01: Pod Basics

---

If this repo is helpful, don't forget to star it! # Lab 01: Pod Basics

## Learning Objectives

- Understand what a Pod is
  - Create Pods (imperative and declarative)
  - List and inspect Pods
  - Delete Pods
- 

## What is a Pod?

### □ Why is Pod Important?

A **Pod** is the fundamental building block of Kubernetes. What a “container” is in Docker, a “Pod” is in Kubernetes.

Concept	Docker	Kubernetes
Smallest unit	Container	Pod
Run command	<code>docker run</code>	<code>kubectl run</code>
Network	Container network	Pod network

## Real-World Example

Think of a website: - **Container 1:** Web server (nginx) - **Container 2:** Log collector (fluentd)  
- **Shared Volume:** Log files

These run in the same Pod because: - They share the same lifecycle - They communicate closely - They share data

## Key Pod Properties




Figure 6: Mermaid Diagram

Property	Description
<b>IP Address</b>	Each Pod gets its own cluster IP
<b>Shared Network</b>	Containers in a Pod talk via localhost
<b>Ephemeral</b>	If a Pod dies, a new one is created (not the same one)
<b>Disposable</b>	Don't create Pods directly, use Deployment!

## Big Picture: Where Does Pod Fit?




Figure 7: Mermaid Diagram

**Important:** In the real world, we don't create Pods directly! We use Deployments (Lab 03). But you can't understand Deployments without understanding Pods.

```
kubectl get pods
```

---

## Exercise 2: Watch Pod Status

**Task:** Watch the Pod transition to Running state.

Solution

```
# Watch mode (Ctrl+C to exit)
kubectl get pods -w
```

### Pod States (Lifecycle):

```
Pending → ContainerCreating → Running
      ↓           ↓           ↓
Queued     Image is       Working!
            being pulled
```

If there's an error: - ImagePullBackOff → Wrong image name - CrashLoopBackOff → Container keeps crashing - Error → Something is wrong

---

## Exercise 3: Inspect Pod Details

**Task:** Get detailed information about my-first-pod.

**What's this for?** This is the most important command for debugging!

Solution

```
kubectl describe pod my-first-pod
```

### Important sections to look at:

Section	Description
<b>Node</b>	Which server is the Pod running on
<b>IP</b>	Pod's cluster-internal IP address
<b>Containers</b>	Container status and restart count
<b>Events</b>	Recent events (for finding errors!)

---

## Exercise 4: Create Pod with YAML

**Task:** Create a pod YAML file with these specs: - Name: redis-pod - Image: redis:alpine - Label: app=cache

**What's this for?** YAML = Infrastructure as Code. Keep all settings in a file, use version control.

Hint - Exam Trick!

You don't need to memorize YAML! Kubernetes gives you a template:

```
kubectl run redis-pod --image=redis:alpine --labels=app=cache --dry-run=client -o yaml
--dry-run=client → Don't actually create, just show YAML -o yaml → Output in YAML format
```

Solution

```

# 1. Generate YAML template
kubectl run redis-pod --image=redis:alpine --labels=app=cache --dry-run=client -o yaml > redis-pod.yaml

# 2. Inspect the file
cat redis-pod.yaml

# 3. Apply
kubectl apply -f redis-pod.yaml

```

### **YAML Explained:**

```

apiVersion: v1           # API version
kind: Pod                # Resource type
metadata:
  name: redis-pod        # Pod name
  labels:
    app: cache            # Label (for filtering)
spec:
  containers:
  - name: redis           # Container name
    image: redis:alpine   # Image to use

```

---

### **Exercise 5: View Pod Logs**

**Task:** View logs from my-first-pod.

**What's this for?** What's happening inside the container? Any errors? Logs will tell you.

Solution

```

# Current logs
kubectl logs my-first-pod

# Follow logs live (Ctrl+C to exit)
kubectl logs -f my-first-pod

# Last 10 lines
kubectl logs --tail=10 my-first-pod

# Previous (crashed) container logs
kubectl logs my-first-pod --previous

```

---

### **Exercise 6: Execute Commands in Pod**

**Task:** Open a shell in my-first-pod and run hostname.

**What's this for?** Debug inside the container, check files, test network.

Solution

```

# Run a single command
kubectl exec my-first-pod -- hostname

# Open interactive shell
kubectl exec -it my-first-pod -- /bin/sh

```

```
# Commands to try inside:  
# ls -la  
# cat /etc/nginx/nginx.conf  
# curl localhost:80  
# exit
```

**What does -it mean?** - -i = interactive (stdin open) - -t = TTY (terminal)

---

### Exercise 7: Filter with Labels

**Task:** List pods with label app=cache.

**What's this for?** Labels are Kubernetes' "search engine". Find what you want among thousands of pods!

Solution

```
# Filter by label  
kubectl get pods -l app=cache  
  
# Show all labels  
kubectl get pods --show-labels  
  
# Add a label  
kubectl label pod my-first-pod env=dev  
  
# Remove a label (- at the end)  
kubectl label pod my-first-pod env-
```

---

### Exercise 8: Delete Pods

**Task:** Delete the pods you created.

Solution

```
# Delete single pod  
kubectl delete pod my-first-pod  
  
# Delete using YAML (deletes what you created)  
kubectl delete -f redis-pod.yaml  
  
# Delete by label  
kubectl delete pods -l app=cache  
  
# Delete all (CAREFUL!)  
kubectl delete pods --all  
  
# Quick delete (for exam)  
kubectl delete pod my-first-pod --force --grace-period=0
```

---

## Exam Practice

Solve these scenarios with a timer! Target: < 2 minutes each

### Scenario 1

Create a pod named `test-pod` using `busybox` image. The pod should run `sleep 3600` command.

Solution

```
kubectl run test-pod --image=busybox --command -- sleep 3600
```

---

### Scenario 2

Create a pod named `webapp` using `nginx:1.21` image with label `tier=frontend`.

Solution

```
kubectl run webapp --image=nginx:1.21 --labels=tier=frontend
```

---

### Scenario 3

Find which node the `webapp` pod is running on.

Solution

```
kubectl get pod webapp -o wide  
# or  
kubectl describe pod webapp | grep Node
```

---

## Common Mistakes

Error	Symptom	Solution
Wrong image name	ImagePullBackOff	Check image name
Wrong command	CrashLoopBackOff	Check with kubectl logs
Port conflict	Error	Use different port
YAML syntax error	error parsing	Check YAML indentation

---

## □ Cleanup

```
kubectl delete pod --all  
rm -f redis-pod.yaml
```

---

## What We Learned

- Pod = Kubernetes' smallest unit
  - kubectl run to create pods
  - kubectl get pods to list
  - kubectl describe for details
  - kubectl logs for logs
  - kubectl exec to run commands inside
  - dry-run=client -o yaml for YAML generation (exam trick!)
  - Label filtering
- 

## Next Steps

You understand Pods. But is a single Pod enough? What if the Pod dies?

Lab 02: Multi-Container Pods - Put multiple containers in a Pod

Lab 03: Deployments - Manage Pods automatically, recreate if they die

---

[Home](#) | [Lab 02: Multi-Container Pods](#) # [Lab 02: Multi-Container Pods](#)

## Learning Objectives

- Understand multi-container pod patterns
  - Implement Sidecar pattern
  - Use Init Containers
  - Communication between containers
- 

## Multi-Container Patterns




Figure 8: Mermaid Diagram

Pattern	Use Case
<b>Sidecar</b>	Log collection, sync, proxy
<b>Ambassador</b>	Connection to external services
<b>Adapter</b>	Data format conversion

---

## Hands-on Exercises

### Exercise 1: Two-Container Pod

**Task:** Create a pod with these specs: - Pod name: two-containers - Container 1: nginx (name: web) - Container 2: busybox (name: sidecar), runs sleep 3600

## Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: two-containers
spec:
  containers:
    - name: web
      image: nginx
      ports:
        - containerPort: 80
    - name: sidecar
      image: busybox
      command: ["sleep", "3600"]

kubectl apply -f two-containers.yaml
kubectl get pods two-containers
```

---

## Exercise 2: Access Specific Container

**Task:** Connect to the `sidecar` container in `two-containers` pod.

### Solution

```
# Exec into specific container
kubectl exec -it two-containers -c sidecar -- /bin/sh

# Logs from specific container
kubectl logs two-containers -c web
kubectl logs two-containers -c sidecar
```

---

## Exercise 3: Sidecar with Shared Volume

**Task:** Main container writes logs, sidecar reads them.

### Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: sidecar-pod
spec:
  containers:
    - name: web
      image: nginx
      volumeMounts:
        - name: logs
          mountPath: /var/log/nginx

    - name: log-reader
      image: busybox
      command: ["sh", "-c", "tail -f /logs/access.log 2>/dev/null || sleep 3600"]
      volumeMounts:
```




Figure 9: Mermaid Diagram

```

- name: logs
  mountPath: /logs

volumes:
- name: logs
  emptyDir: {}

```

---

#### Exercise 4: Init Container




Figure 10: Mermaid Diagram

**Task:** Add an init container that prepares a file before main container starts.

Solution

```

apiVersion: v1
kind: Pod
metadata:
  name: init-pod
spec:
  initContainers:
  - name: init-download
    image: busybox
    command: ['sh', '-c', 'echo "Ready: $(date)" > /work/status.txt']

```

```

volumeMounts:
- name: workdir
  mountPath: /work

containers:
- name: main-app
  image: busybox
  command: ['sh', '-c', 'cat /work/status.txt && sleep 3600']
  volumeMounts:
- name: workdir
  mountPath: /work

volumes:
- name: workdir
  emptyDir: {}

```

---

## Exercise 5: Multiple Init Containers

**Task:** Create two init containers that run sequentially.

Solution

```

apiVersion: v1
kind: Pod
metadata:
  name: multi-init-pod
spec:
  initContainers:
  - name: init-1
    image: busybox
    command: ['sh', '-c', 'echo "Step 1" && sleep 2']

  - name: init-2
    image: busybox
    command: ['sh', '-c', 'echo "Step 2" && sleep 2']

  containers:
  - name: main
    image: nginx

```

Init containers run sequentially. One must finish before the next starts.

---

## Exercise 6: Inter-Container Networking

**Task:** Test that containers in the same pod can communicate via localhost.

Solution

```

apiVersion: v1
kind: Pod
metadata:
  name: network-test
spec:

```

```

containers:
- name: web
  image: nginx
- name: tester
  image: curlimages/curl
  command: ["sleep", "3600"]

kubectl apply -f network-test.yaml

# Access web from tester (localhost)
kubectl exec network-test -c tester -- curl -s localhost:80
Containers in the same pod communicate via localhost!

```

---

## Exam Practice

### Scenario 1

Create a pod named log-app: - Container 1: nginx (name: app) - Container 2: busybox (name: logger), runs sleep 3600

Solution

```

apiVersion: v1
kind: Pod
metadata:
  name: log-app
spec:
  containers:
    - name: app
      image: nginx
    - name: logger
      image: busybox
      command: ["sleep", "3600"]

```

---

### Scenario 2

Create a pod with init container. Init container runs wget -O /data/index.html http://info.cern.ch. Main container is nginx and serves this file.

Solution

```

apiVersion: v1
kind: Pod
metadata:
  name: web-init
spec:
  initContainers:
    - name: downloader
      image: busybox
      command: ['wget', '-O', '/data/index.html', 'http://info.cern.ch']
  volumeMounts:
    - name: html
      mountPath: /data

```

```
containers:
- name: nginx
  image: nginx
  volumeMounts:
  - name: html
    mountPath: /usr/share/nginx/html

volumes:
- name: html
  emptyDir: {}
```

---

## □ Cleanup

```
kubectl delete pod two-containers sidecar-pod init-pod multi-init-pod network-test --ignore-not-found
```

---

## What We Learned

- ☒ Multi-container pod creation
  - ☒ Sidecar pattern
  - ☒ Init containers
  - ☒ Shared volumes between containers
  - ☒ -c flag for container selection
  - ☒ Localhost communication between containers
- 

Lab 01: Pod Basics | Lab 03: Deployments # Lab 03: Deployments

## Learning Objectives

- Understand what a Deployment is and why it's used
  - Create and manage Deployments
  - Rolling Update and Rollback
  - Scaling
- 

## What is a Deployment?

Deployment provides:  
- Automatic pod recreation  
- Rolling updates (zero-downtime)  
- Rollback capability  
- Scaling

---

## Hands-on Exercises

### Exercise 1: Create Deployment

**Task:** Create a deployment named `web-deploy` using `nginx` image with 3 replicas.

Solution




Figure 11: Mermaid Diagram

```
kubectl create deployment web-deploy --image=nginx --replicas=3
```

Check:

```
kubectl get deployments
kubectl get replicaset
kubectl get pods
```

---

## Exercise 2: Deployment YAML

**Task:** Create a deployment using a YAML file.

Solution

```
# Generate template
kubectl create deployment api-deploy --image=nginx:1.20 --replicas=2 --dry-run=client -o yaml > api-deploy.yaml

apiVersion: apps/v1
kind: Deployment
metadata:
  name: api-deploy
spec:
  replicas: 2
  selector:
    matchLabels:
      app: api-deploy
  template:
    metadata:
      labels:
        app: api-deploy
```

```
spec:  
  containers:  
    - name: nginx  
      image: nginx:1.20  
kubectl apply -f api-deploy.yaml
```

---

### Exercise 3: Scaling

**Task:** Scale web-deploy to 5 replicas.

Solution

```
kubectl scale deployment web-deploy --replicas=5
```

```
# Check  
kubectl get deployment web-deploy  
kubectl get pods -l app=web-deploy
```

---

### Exercise 4: Rolling Update

**Task:** Update web-deploy image to nginx:1.21.

Solution

```
# Update image  
kubectl set image deployment/web-deploy nginx=nginx:1.21  
  
# Watch rollout status  
kubectl rollout status deployment/web-deploy
```

---

### Exercise 5: Rollout History

**Task:** View deployment update history.

Solution

```
# View history  
kubectl rollout history deployment/web-deploy  
  
# Specific revision details  
kubectl rollout history deployment/web-deploy --revision=1
```

---

### Exercise 6: Rollback

**Task:** Rollback deployment to previous version.

Solution

```
# Rollback to previous  
kubectl rollout undo deployment/web-deploy
```




Figure 12: Mermaid Diagram




Figure 13: Mermaid Diagram

```
# Rollback to specific revision
kubectl rollout undo deployment/web-deploy --to-revision=1

# Check status
kubectl rollout status deployment/web-deploy
```

---

## Exercise 7: Deployment Strategies

**RollingUpdate** (default):

```
spec:
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1          # How many extra pods
      maxUnavailable: 1   # How many pods can be unavailable
```

**Recreate** (delete all, create new):

```
spec:
  strategy:
    type: Recreate
```

---

## Exercise 8: Pause and Resume

**Task:** Pause rollout, make multiple changes, then resume.

Solution

```
# Pause
kubectl rollout pause deployment/web-deploy

# Make changes (rollout won't start)
kubectl set image deployment/web-deploy nginx=nginx:1.22
kubectl set resources deployment/web-deploy -c nginx --limits=memory=256Mi

# Resume (single rollout)
kubectl rollout resume deployment/web-deploy
```

---

## Exam Practice

### Scenario 1

Create a deployment named `frontend` using `httpd:2.4` image with 4 replicas.

Solution

```
kubectl create deployment frontend --image=httpd:2.4 --replicas=4
```

---

## Scenario 2

Update frontend image to httpd:alpine. Then rollback to revision 1.

Solution

```
kubectl set image deployment/frontend httpd=httpd:alpine
kubectl rollout status deployment/frontend
kubectl rollout undo deployment/frontend --to-revision=1
```

---

## Scenario 3

Create a redis deployment named backend with 2 replicas. Then scale to 6 replicas.

Solution

```
kubectl create deployment backend --image=redis --replicas=2
kubectl scale deployment backend --replicas=6
```

---

## □ Cleanup

```
kubectl delete deployment --all
```

---

## What We Learned

- Deployment creation
  - ReplicaSet relationship
  - Scaling
  - Rolling update
  - Rollback
  - Deployment strategies
- 

Lab 02 | Lab 04: Services # Lab 04: Services

## Learning Objectives

- Understand what a Service is
  - Service types (ClusterIP, NodePort, LoadBalancer)
  - Service discovery and DNS
  - Expose deployments
- 

## What is a Service?

Type	Access	Use Case
<b>ClusterIP</b>	Internal only	Pod-to-pod communication
<b>NodePort</b>	NodeIP:Port	Development, testing
<b>LoadBalancer</b>	External IP	Production (cloud)

Type	Access	Use Case

## Hands-on Exercises

### Exercise 1: Create Deployment for Testing

```
kubectl create deployment web --image=nginx --replicas=3
```

---

### Exercise 2: ClusterIP Service

**Task:** Expose `web` deployment with ClusterIP service.

Solution

```
kubectl expose deployment web --port=80 --type=ClusterIP
```

Check:

```
kubectl get svc web
kubectl describe svc web
kubectl get endpoints web
```

---

### Exercise 3: NodePort Service

**Task:** Create a NodePort service for `web` deployment.

Solution

```
kubectl expose deployment web --port=80 --type=NodePort --name=web-nodeport
kubectl get svc web-nodeport
# Note the NodePort (e.g., 32000)

# Access from outside
curl http://<NODE_IP>:<NODE_PORT>
```

---

### Exercise 4: Service YAML

Solution

```
apiVersion: v1
kind: Service
metadata:
  name: web-svc
spec:
  type: ClusterIP
  selector:
    app: web      # Must match pod labels
  ports:
```




Figure 14: Mermaid Diagram

```
- port: 80      # Service port
  targetPort: 80 # Container port
```

---

### Exercise 5: Service DNS

**Task:** Test service DNS resolution.




Figure 15: Mermaid Diagram

Solution

```
# Create a test pod
kubectl run test-dns --image=busybox --rm -it --restart=Never -- nslookup web

# Full DNS name
# <service>.〈namespace〉.svc.cluster.local
# web.default.svc.cluster.local
```

---

### Exercise 6: Endpoints

**Task:** View service endpoints.

Solution

```
kubectl get endpoints web

# Shows pod IPs that match the selector
# When pods scale, endpoints update automatically
```

---

### Exercise 7: LoadBalancer (K3s)

K3s has built-in LoadBalancer support (ServiceLB).

Solution

```
kubectl expose deployment web --port=80 --type=LoadBalancer --name=web-lb

kubectl get svc web-lb
# EXTERNAL-IP will show node IP in K3s
```

---

## **Exercise 8: Multi-Port Service**

**Task:** Create a service with multiple ports.

Solution

```
apiVersion: v1
kind: Service
metadata:
  name: multi-port-svc
spec:
  selector:
    app: myapp
  ports:
  - name: http
    port: 80
    targetPort: 80
  - name: https
    port: 443
    targetPort: 443
```

---

## **Exam Practice**

### **Scenario 1**

Create a ClusterIP service named backend-svc for pods with label app=backend, port 8080.

Solution

```
apiVersion: v1
kind: Service
metadata:
  name: backend-svc
spec:
  selector:
    app: backend
  ports:
  - port: 8080
```

---

### **Scenario 2**

Expose deployment api on NodePort 30080.

Solution

```
apiVersion: v1
kind: Service
metadata:
  name: api-svc
spec:
  type: NodePort
  selector:
    app: api
```

```
ports:  
- port: 80  
  nodePort: 30080
```

---

## □ Cleanup

```
kubectl delete deployment web  
kubectl delete svc web web-nodeport web-lb multi-port-svc --ignore-not-found
```

---

## What We Learned

- ☒ Service types (ClusterIP, NodePort, LoadBalancer)
  - ☒ kubectl expose command
  - ☒ Service selectors and endpoints
  - ☒ DNS resolution in Kubernetes
  - ☒ Multi-port services
- 

Lab 03 | Lab 05: ConfigMaps & Secrets # Lab 05: ConfigMaps & Secrets

## Learning Objectives

- Create and use ConfigMaps
  - Create and use Secrets
  - Environment variables from ConfigMaps/Secrets
  - Mount as volumes
- 

## ConfigMap vs Secret

Feature	ConfigMap	Secret
<b>Data Type</b>	Plain text	Base64 encoded
<b>Use Case</b>	Config files, env vars	Passwords, tokens, keys
<b>Visibility</b>	Visible in <code>kubectl get</code>	Hidden by default

---

## ConfigMap Exercises

### Exercise 1: Create ConfigMap (literal)

**Task:** Create a ConfigMap named `app-config` with key `APP_ENV=production`.

Solution

```
kubectl create configmap app-config --from-literal=APP_ENV=production --from-literal=LOG_LEVEL=info
```

Check:




Figure 16: Mermaid Diagram

---

```
kubectl get cm app-config
kubectl describe cm app-config
```

---

### **Exercise 2: Create ConfigMap (from file)**

Solution

```
# Create a config file
echo "database_host=db.example.com
database_port=5432" > config.txt

kubectl create configmap file-config --from-file=config.txt
```

---

### **Exercise 3: ConfigMap as Environment Variables**

**Task:** Use ConfigMap values as environment variables in a pod.

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: cm-env-pod
spec:
  containers:
  - name: app
    image: busybox
    command: ["sh", "-c", "echo $APP_ENV && sleep 3600"]
    envFrom:
```

```
- configMapRef:  
  name: app-config
```

Or specific keys:

```
env:  
- name: MY_APP_ENV  
  valueFrom:  
    configMapKeyRef:  
      name: app-config  
      key: APP_ENV
```

---

#### Exercise 4: ConfigMap as Volume

**Task:** Mount ConfigMap as a file.

Solution

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: cm-vol-pod  
spec:  
  containers:  
    - name: app  
      image: busybox  
      command: ["sh", "-c", "cat /config/config.txt && sleep 3600"]  
      volumeMounts:  
        - name: config-volume  
          mountPath: /config  
    volumes:  
      - name: config-volume  
        configMap:  
          name: file-config
```

---

### Secret Exercises

#### Exercise 5: Create Secret

**Task:** Create a Secret named db-secret with DB\_PASSWORD=mysecret123.

Solution

```
kubectl create secret generic db-secret --from-literal=DB_PASSWORD=mysecret123
```

Check:

```
kubectl get secret db-secret  
kubectl describe secret db-secret # Values hidden  
kubectl get secret db-secret -o yaml # Base64 encoded
```

---

## Exercise 6: Decode Secret

Solution

```
# Get base64 value
kubectl get secret db-secret -o jsonpath='{.data.DB_PASSWORD}'

# Decode
kubectl get secret db-secret -o jsonpath='{.data.DB_PASSWORD}' | base64 -d
```

---

## Exercise 7: Secret as Environment Variable

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: secret-env-pod
spec:
  containers:
    - name: app
      image: busybox
      command: ["sh", "-c", "echo $DB_PASSWORD && sleep 3600"]
      env:
        - name: DB_PASSWORD
          valueFrom:
            secretKeyRef:
              name: db-secret
              key: DB_PASSWORD
```

---

## Exercise 8: Secret as Volume

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: secret-vol-pod
spec:
  containers:
    - name: app
      image: busybox
      command: ["sh", "-c", "cat /secrets/DB_PASSWORD && sleep 3600"]
      volumeMounts:
        - name: secret-volume
          mountPath: /secrets
          readOnly: true
  volumes:
    - name: secret-volume
      secret:
        secretName: db-secret
```

---

## Exam Practice

### Scenario 1

Create ConfigMap webapp-config with THEME=dark and CACHE\_TTL=3600.

Solution

```
kubectl create cm webapp-config --from-literal=THEME=dark --from-literal=CACHE_TTL=3600
```

---

### Scenario 2

Create a pod that uses webapp-config as environment variables.

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: webapp
spec:
  containers:
    - name: app
      image: nginx
      envFrom:
        - configMapRef:
            name: webapp-config
```

---

### □ Cleanup

```
kubectl delete pod --all
kubectl delete cm app-config file-config webapp-config --ignore-not-found
kubectl delete secret db-secret --ignore-not-found
rm -f config.txt
```

---

## What We Learned

- Create ConfigMaps (literal, file)
- Create Secrets
- Use as environment variables
- Mount as volumes
- Base64 encoding/decoding

---

Lab 04 | Lab 06: Volumes # Lab 06: Volumes & PersistentVolumes

## Learning Objectives

- Understand volume types
- Use emptyDir and hostPath
- Create PersistentVolume and PersistentVolumeClaim

- K3s local-path provisioner
- 

## Volume Types




Figure 17: Mermaid Diagram

Type	Lifetime	Use Case
<b>emptyDir</b>	Pod lifetime	Temp files, cache
<b>hostPath</b>	Node lifetime	Testing only
<b>PV/PVC</b>	Independent	Production data

---

## Hands-on Exercises

### Exercise 1: emptyDir

**Task:** Create a pod with emptyDir volume shared between containers.

Solution

```

apiVersion: v1
kind: Pod
metadata:
  name: emptydir-pod
spec:
  containers:
    - name: writer
      image: busybox
      command: ["sh", "-c", "echo 'Hello' > /data/file.txt && sleep 3600"]
      volumeMounts:
        - name: shared-data
          mountPath: /data

    - name: reader
      image: busybox
      command: ["sh", "-c", "cat /data/file.txt && sleep 3600"]
      volumeMounts:
        - name: shared-data
          mountPath: /data

  volumes:
    - name: shared-data
      emptyDir: {}

```

---

## Exercise 2: hostPath

**Task:** Mount node's /tmp directory into a pod.

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: hostpath-pod
spec:
  containers:
    - name: app
      image: busybox
      command: ["sleep", "3600"]
      volumeMounts:
        - name: host-volume
          mountPath: /host-data
  volumes:
    - name: host-volume
      hostPath:
        path: /tmp
        type: Directory
```

**Warning:** hostPath is not recommended for production!

---

## Exercise 3: PersistentVolume (Static)

**Task:** Create a PersistentVolume.

Solution

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: /tmp/my-pv
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 500Mi
```

```
kubectl apply -f pv-pvc.yaml  
kubectl get pv  
kubectl get pvc
```

---

### Exercise 4: Use PVC in Pod

Solution

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: pvc-pod  
spec:  
  containers:  
    - name: app  
      image: nginx  
      volumeMounts:  
        - name: data  
          mountPath: /usr/share/nginx/html  
  volumes:  
    - name: data  
      persistentVolumeClaim:  
        claimName: my-pvc
```

---

### Exercise 5: K3s Dynamic Provisioning

K3s includes local-path provisioner for automatic PV creation.

Solution

```
apiVersion: v1  
kind: PersistentVolumeClaim  
metadata:  
  name: dynamic-pvc  
spec:  
  accessModes:  
    - ReadWriteOnce  
  storageClassName: local-path # K3s default  
  resources:  
    requests:  
      storage: 1Gi  
  
kubectl apply -f dynamic-pvc.yaml  
kubectl get pvc  
# Status will be Bound automatically!
```

---

### Exercise 6: Access Modes

Mode	Description
<b>ReadWriteOnce (RWO)</b>	Single node read/write
<b>ReadOnlyMany (ROX)</b>	Multiple nodes read-only
<b>ReadWriteMany (RWX)</b>	Multiple nodes read/write

## Exam Practice

### Scenario 1

Create a pod with emptyDir volume mounted at /cache.

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: cache-pod
spec:
  containers:
    - name: app
      image: nginx
      volumeMounts:
        - name: cache-vol
          mountPath: /cache
  volumes:
    - name: cache-vol
      emptyDir: {}
```

### Scenario 2

Create PVC named data-pvc requesting 2Gi storage.

Solution

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: data-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 2Gi
```

## □ Cleanup

```
kubectl delete pod emptydir-pod hostpath-pod pvc-pod cache-pod --ignore-not-found
kubectl delete pvc my-pvc dynamic-pvc data-pvc --ignore-not-found
kubectl delete pv my-pv --ignore-not-found
```

---

## What We Learned

- emptyDir for temporary storage
  - hostPath for node storage
  - PersistentVolume and PersistentVolumeClaim
  - Dynamic provisioning with StorageClass
  - K3s local-path provisioner
- 

Lab 05 | Lab 07: Jobs & CronJobs # Lab 07: Jobs & CronJobs

## Learning Objectives

- Understand Jobs for one-time tasks
  - Understand CronJobs for scheduled tasks
  - Parallelism and completions
  - Cron syntax
- 

## Job vs CronJob




Figure 18: Mermaid Diagram

---

Resource	Use Case
<b>Job</b>	Batch processing, one-time tasks
<b>CronJob</b>	Scheduled backups, reports

---

## Job Exercises

### Exercise 1: Simple Job

**Task:** Create a Job that prints “Hello CKAD” and exits.

## Solution

```
kubectl create job hello-job --image=busybox -- echo "Hello CKAD"
```

Check:

```
kubectl get jobs  
kubectl get pods  
kubectl logs <pod-name>
```

---

## Exercise 2: Job YAML

### Solution

```
apiVersion: batch/v1  
kind: Job  
metadata:  
  name: pi-job  
spec:  
  template:  
    spec:  
      containers:  
      - name: pi  
        image: perl  
        command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(100)"]  
        restartPolicy: Never  
    backoffLimit: 4
```

---

## Exercise 3: Job with Completions

**Task:** Create a Job that runs 5 times.

### Solution

```
apiVersion: batch/v1  
kind: Job  
metadata:  
  name: multi-job  
spec:  
  completions: 5      # Run 5 times  
  parallelism: 2      # 2 at a time  
  template:  
    spec:  
      containers:  
      - name: worker  
        image: busybox  
        command: ["sh", "-c", "echo Processing && sleep 5"]  
    restartPolicy: Never
```

---

## Exercise 4: Job Cleanup

### Solution

```
spec:  
  ttlSecondsAfterFinished: 60  # Auto-delete after 60s
```

---

## CronJob Exercises

### Exercise 5: Simple CronJob

**Task:** Create a CronJob that runs every minute.

Solution

```
kubectl create cronjob minute-cron --image=busybox --schedule="* * * * *" -- date
```

Check:

```
kubectl get cronjobs  
kubectl get jobs  # A new job each minute
```

---

### Exercise 6: Cron Syntax

```
minute (0 - 59)  
hour (0 - 23)  
day of month (1 - 31)  
month (1 - 12)  
day of week (0 - 6)
```

```
* * * * *
```

Schedule	Meaning
* * * * *	Every minute
0 * * * *	Every hour
0 0 * * *	Every day at midnight
0 0 * * 0	Every Sunday
*/5 * * * *	Every 5 minutes

---

### Exercise 7: CronJob YAML

Solution

```
apiVersion: batch/v1  
kind: CronJob  
metadata:  
  name: backup-cron  
spec:  
  schedule: "0 2 * * *"  # 2 AM daily  
  jobTemplate:  
    spec:  
      template:  
        spec:  
          containers:
```

```
- name: backup
  image: busybox
  command: ["sh", "-c", "echo Backup at $(date)"]
  restartPolicy: OnFailure
```

---

## Exercise 8: Concurrency Policy

Policy	Behavior
<b>Allow</b>	Multiple jobs can run
<b>Forbid</b>	Skip if previous running
<b>Replace</b>	Replace previous job

Solution

```
spec:
  concurrencyPolicy: Forbid
```

---

## Exam Practice

### Scenario 1

Create a Job named count-job that counts from 1 to 10.

Solution

```
kubectl create job count-job --image=busybox -- sh -c "for i in $(seq 1 10); do echo $i; done"
```

---

### Scenario 2

Create a CronJob named report-cron that runs every 5 minutes.

Solution

```
kubectl create cronjob report-cron --image=busybox --schedule="*/5 * * * *" -- echo "Report generated"
```

---

## □ Cleanup

```
kubectl delete job --all
kubectl delete cronjob --all
```

---

## What We Learned

- Create Jobs
- Completions and parallelism
- Create CronJobs
- Cron syntax

- ☒ Concurrency policies
- 

Lab 06 | Lab 08: Probes # Lab 08: Liveness & Readiness Probes

## Learning Objectives

- Understand probe types
  - Implement Liveness probes
  - Implement Readiness probes
  - Startup probes
- 

## Probe Types




Figure 19: Mermaid Diagram

Probe	Question	Action on Failure
<b>Liveness</b>	Is the app alive?	Restart container
<b>Readiness</b>	Is the app ready?	Remove from Service endpoints
<b>Startup</b>	Has the app started?	Block other probes

---

## Hands-on Exercises

### Exercise 1: HTTP Liveness Probe

**Task:** Create a pod with HTTP liveness probe.

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: liveness-http
spec:
  containers:
  - name: app
    image: nginx
    livenessProbe:
      httpGet:
        path: /
        port: 80
      initialDelaySeconds: 5
      periodSeconds: 10
```

---

### Exercise 2: TCP Liveness Probe

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: liveness-tcp
spec:
  containers:
  - name: app
    image: redis
    livenessProbe:
      tcpSocket:
        port: 6379
      initialDelaySeconds: 5
      periodSeconds: 10
```

---

### Exercise 3: Exec Liveness Probe

**Task:** Run a command to check health.

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: liveness-exec
spec:
  containers:
  - name: app
```

```
image: busybox
command: ["sh", "-c", "touch /tmp/healthy && sleep 3600"]
livenessProbe:
  exec:
    command:
      - cat
      - /tmp/healthy
  initialDelaySeconds: 5
  periodSeconds: 5
```

---

## Exercise 4: Readiness Probe

**Task:** Pod only receives traffic when ready.

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: readiness-pod
spec:
  containers:
    - name: app
      image: nginx
      readinessProbe:
        httpGet:
          path: /
          port: 80
      initialDelaySeconds: 5
      periodSeconds: 5
```

---

## Exercise 5: Both Probes

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: full-probe
spec:
  containers:
    - name: app
      image: nginx
      livenessProbe:
        httpGet:
          path: /healthz
          port: 80
      initialDelaySeconds: 10
      periodSeconds: 15
      readinessProbe:
        httpGet:
          path: /ready
```

```
  port: 80
  initialDelaySeconds: 5
  periodSeconds: 5
```

---

## Exercise 6: Probe Parameters

Parameter	Description	Default
initialDelaySeconds	Wait before first probe	0
periodSeconds	How often to probe	10
timeoutSeconds	Timeout for probe	1
successThreshold	Min successes for success	1
failureThreshold	Min failures for failure	3

---

## Exercise 7: Startup Probe

For slow-starting applications.

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: slow-start
spec:
  containers:
  - name: app
    image: nginx
    startupProbe:
      httpGet:
        path: /
        port: 80
      failureThreshold: 30
      periodSeconds: 10
    livenessProbe:
      httpGet:
        path: /
        port: 80
      periodSeconds: 10
```

---

## Exam Practice

### Scenario 1

Create a pod `web-health` with HTTP liveness probe on port 80, path `/health`.

Solution

```
apiVersion: v1
kind: Pod
```

```
metadata:  
  name: web-health  
spec:  
  containers:  
    - name: web  
      image: nginx  
      livenessProbe:  
        httpGet:  
          path: /health  
          port: 80
```

---

## Scenario 2

Add readiness probe to existing pod checking TCP port 3306.

Solution

```
readinessProbe:  
  tcpSocket:  
    port: 3306  
  initialDelaySeconds: 5  
  periodSeconds: 10
```

---

## □ Cleanup

```
kubectl delete pod liveness-http liveness-tcp liveness-exec readiness-pod full-probe slow-start web-hea
```

---

## What We Learned

- ☒ Liveness vs Readiness vs Startup probes
  - ☒ HTTP, TCP, and Exec probe types
  - ☒ Probe parameters
  - ☒ When to use each probe type
- 

Lab 07 | Lab 09: Resource Limits # Lab 09: Resource Limits & Requests

## Learning Objectives

- Understand CPU and Memory requests/limits
  - Configure resource constraints
  - LimitRange and ResourceQuota
- 

## Requests vs Limits




Figure 20: Mermaid Diagram

Concept	Description
<b>Request</b>	Minimum guaranteed resources
<b>Limit</b>	Maximum allowed resources

## Hands-on Exercises

### Exercise 1: Set Resource Requests and Limits

Solution

```

apiVersion: v1
kind: Pod
metadata:
  name: resource-pod
spec:
  containers:
  - name: app
    image: nginx
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"      # 0.25 CPU
      limits:
        memory: "128Mi"
        cpu: "500m"      # 0.5 CPU

```

## Exercise 2: CPU Units

Value	Meaning
1	1 full CPU
500m	0.5 CPU (500 millicores)
100m	0.1 CPU

---

## Exercise 3: Memory Units

Value	Meaning
128Mi	128 Mebibytes
1Gi	1 Gibibyte
500M	500 Megabytes

---

## Exercise 4: LimitRange

**Task:** Set default limits for a namespace.

Solution

```
apiVersion: v1
kind: LimitRange
metadata:
  name: default-limits
spec:
  limits:
  - default:          # Default limits
    cpu: "500m"
    memory: "256Mi"
  defaultRequest:    # Default requests
    cpu: "100m"
    memory: "64Mi"
  max:               # Maximum allowed
    cpu: "1"
    memory: "1Gi"
  min:               # Minimum allowed
    cpu: "50m"
    memory: "32Mi"
  type: Container
```

---

## Exercise 5: ResourceQuota

**Task:** Limit total resources in a namespace.

Solution

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: ns-quota
spec:
  hard:
    requests.cpu: "2"
    requests.memory: "2Gi"
    limits.cpu: "4"
    limits.memory: "4Gi"
    pods: "10"

kubectl apply -f quota.yaml
kubectl describe quota ns-quota
```

---

## Exercise 6: View Resource Usage

Solution

```
# Pod resources
kubectl top pods

# Node resources
kubectl top nodes

# Describe pod for limits
kubectl describe pod <pod-name> | grep -A5 "Limits"
```

---

## Exam Practice

### Scenario 1

Create a pod with 100m CPU request and 200m CPU limit.

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: cpu-pod
spec:
  containers:
  - name: app
    image: nginx
    resources:
      requests:
        cpu: "100m"
      limits:
        cpu: "200m"
```

---

## Scenario 2

Create ResourceQuota limiting namespace to 5 pods.

Solution

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: pod-quota
spec:
  hard:
    pods: "5"
```

---

## □ Cleanup

```
kubectl delete pod resource-pod cpu-pod --ignore-not-found
kubectl delete limitrange default-limits --ignore-not-found
kubectl delete quota ns-quota pod-quota --ignore-not-found
```

---

## What We Learned

- CPU and Memory units
  - Requests vs Limits
  - LimitRange for defaults
  - ResourceQuota for namespace limits
- 

Lab 08 | Lab 10: Network Policies # Lab 10: Network Policies

## Learning Objectives

- Understand NetworkPolicy
  - Ingress and Egress rules
  - Allow/Deny traffic between pods
- 

## What is NetworkPolicy?

Concept	Description
<b>NetworkPolicy</b>	Firewall rules for pods
<b>Ingress</b>	Incoming traffic rules
<b>Egress</b>	Outgoing traffic rules

**K3s Note:** Default Flannel CNI has limited NetworkPolicy support. Consider Calico for full support.

---




Figure 21: Mermaid Diagram

## Hands-on Exercises

### Exercise 1: Default Deny All

**Task:** Block all traffic to pods.

Solution

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny-all
spec:
  podSelector: {}      # Apply to all pods
  policyTypes:
  - Ingress
  - Egress

```

---

### Exercise 2: Allow Specific Pod

**Task:** Allow traffic from frontend to backend.

Solution

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: backend-allow-frontend
spec:
  podSelector:
    matchLabels:
      app: backend
  policyTypes:
  - Ingress
  ingress:
  - from:
    - podSelector:
      matchLabels:
        app: frontend
  ports:

```

```
- protocol: TCP
  port: 80
```

---

### Exercise 3: Allow Namespace

**Task:** Allow traffic from specific namespace.

Solution

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-namespace
spec:
  podSelector:
    matchLabels:
      app: api
  ingress:
  - from:
    - namespaceSelector:
      matchLabels:
        name: production
```

---

### Exercise 4: Allow DNS (Egress)

**Task:** Allow DNS resolution.

Solution

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-dns
spec:
  podSelector: {}
  policyTypes:
  - Egress
  egress:
  - to:
    - namespaceSelector: {}
  ports:
  - protocol: UDP
    port: 53
```

---

### Exercise 5: Combined Rules

Solution

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
```

```
name: db-policy
spec:
  podSelector:
    matchLabels:
      app: db
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        podSelector:
          matchLabels:
            app: backend
        ports:
          - port: 5432
  egress:
    - to:
        podSelector:
          matchLabels:
            app: backend
```

---

## Exam Practice

### Scenario 1

Create NetworkPolicy that allows ingress to pods labeled `app=web` only from pods labeled `app=api` on port 80.

Solution

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: web-allow-api
spec:
  podSelector:
    matchLabels:
      app: web
  policyTypes:
    - Ingress
  ingress:
    - from:
        podSelector:
          matchLabels:
            app: api
      ports:
        - port: 80
```

---

### Scenario 2

Create default deny egress policy for all pods.

## Solution

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-egress
spec:
  podSelector: {}
  policyTypes:
    - Egress
```

---

## □ Cleanup

```
kubectl delete networkpolicy --all
```

---

## What We Learned

- ☒ NetworkPolicy basics
  - ☒ Ingress and Egress rules
  - ☒ Pod and namespace selectors
  - ☒ Default deny patterns
- 

Lab 09 | Lab 11: Security # Lab 11: Security Context & Service Accounts

## Learning Objectives

- Understand and apply SecurityContext
  - Create and use ServiceAccounts
  - Pod security settings
- 

## What is SecurityContext?

---

## SecurityContext Exercises

### Exercise 1: runAsUser

**Task:** Create a pod running as non-root user.

## Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: security-pod
spec:
  securityContext:
    runAsUser: 1000
```




Figure 22: Mermaid Diagram

```

runAsGroup: 3000
fsGroup: 2000
containers:
- name: app
  image: busybox
  command: ["sh", "-c", "id && sleep 3600"]

kubectl apply -f security-pod.yaml
kubectl logs security-pod
# uid=1000 gid=3000 groups=2000

```

---

### Exercise 2: readOnlyRootFilesystem

Solution

```

apiVersion: v1
kind: Pod
metadata:
  name: readonly-pod
spec:
  containers:
  - name: app
    image: busybox
    command: ["sh", "-c", "sleep 3600"]
    securityContext:
      readOnlyRootFilesystem: true
  volumeMounts:
  - name: tmp
    mountPath: /tmp
  volumes:
  - name: tmp
    emptyDir: {}

```

---

### **Exercise 3: Capabilities**

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: cap-pod
spec:
  containers:
  - name: app
    image: busybox
    command: ["sleep", "3600"]
    securityContext:
      capabilities:
        add: ["NET_ADMIN", "SYS_TIME"]
        drop: ["ALL"]
```

---

### **Exercise 4: allowPrivilegeEscalation**

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: no-escalate-pod
spec:
  containers:
  - name: app
    image: busybox
    command: ["sleep", "3600"]
    securityContext:
      allowPrivilegeEscalation: false
      runAsNonRoot: true
      runAsUser: 1000
```

---

## **ServiceAccount Exercises**

### **Exercise 5: Create ServiceAccount**

Solution

```
kubectl create serviceaccount my-sa
```

Check:

```
kubectl get sa
kubectl describe sa my-sa
```

---

### **Exercise 6: Use ServiceAccount in Pod**

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: sa-pod
spec:
  serviceAccountName: my-sa
  containers:
    - name: app
      image: busybox
      command: ["sleep", "3600"]
```

---

## Exercise 7: Disable Token Mount

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: no-token-pod
spec:
  serviceAccountName: my-sa
  automountServiceAccountToken: false
  containers:
    - name: app
      image: busybox
      command: ["sleep", "3600"]
```

---

## Exam Practice

### Scenario 1

Create a pod running as user 1000, group 3000, with readOnlyRootFilesystem.

Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: secure-pod
spec:
  securityContext:
    runAsUser: 1000
    runAsGroup: 3000
  containers:
    - name: app
      image: busybox
      command: ["sleep", "3600"]
      securityContext:
        readOnlyRootFilesystem: true
```

---

## Scenario 2

Create ServiceAccount backend-sa and use it in a pod named backend.

Solution

```
kubectl create sa backend-sa
```

```
apiVersion: v1
kind: Pod
metadata:
  name: backend
spec:
  serviceAccountName: backend-sa
  containers:
  - name: app
    image: nginx
```

---

## □ Cleanup

```
kubectl delete pod --all
kubectl delete sa my-sa backend-sa --ignore-not-found
```

---

## What We Learned

- Pod and Container SecurityContext
  - runAsUser, runAsGroup, fsGroup
  - readOnlyRootFilesystem
  - Capabilities (add/drop)
  - ServiceAccount creation and usage
- 

Lab 10 | Lab 12: Ingress # Lab 12: Ingress (K3s Traefik)

## Learning Objectives

- Understand Ingress
  - K3s Traefik Ingress Controller
  - Path-based and Host-based routing
  - TLS configuration
- 

## What is Ingress?

K3s comes with **Traefik** Ingress Controller by default!

---




Figure 23: Mermaid Diagram

## Hands-on Exercises

### Setup: Create Test Apps

```
kubectl create deployment app1 --image=nginx --port=80
kubectl expose deployment app1 --port=80
```

```
kubectl create deployment app2 --image=httpd --port=80
kubectl expose deployment app2 --port=80
```

---

### Exercise 1: Simple Ingress

**Task:** Create path-based routing.

Solution

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: simple-ingress
spec:
  rules:
  - http:
      paths:
      - path: /app1
        pathType: Prefix
        backend:
          service:
            name: app1
            port:
              number: 80
      - path: /app2
        pathType: Prefix
        backend:
```

```
service:  
  name: app2  
  port:  
    number: 80
```

---

## Exercise 2: Host-Based Routing

Solution

```
apiVersion: networking.k8s.io/v1  
kind: Ingress  
metadata:  
  name: host-ingress  
spec:  
  rules:  
  - host: app1.local  
    http:  
      paths:  
      - path: /  
        pathType: Prefix  
        backend:  
          service:  
            name: app1  
            port:  
              number: 80  
  - host: app2.local  
    http:  
      paths:  
      - path: /  
        pathType: Prefix  
        backend:  
          service:  
            name: app2  
            port:  
              number: 80
```

Test:

```
curl -H "Host: app1.local" http://<NODE_IP>
```

---

## Exercise 3: pathType Options

pathType	Description
Prefix	Prefix match: /api → /api, /api/v1
Exact	Exact match: /api → only /api

---

## Exercise 4: Default Backend

Solution

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: default-ingress
spec:
  defaultBackend:
    service:
      name: app1
      port:
        number: 80
  rules:
  - http:
      paths:
      - path: /special
        pathType: Prefix
      backend:
        service:
          name: app2
          port:
            number: 80

```

---

## Exercise 5: TLS Ingress

Solution

Create TLS secret:

```
openssl req -x509 -nodes -days 365 -newkey rsa:2048 \
  -keyout tls.key -out tls.crt -subj "/CN=myapp.local"
```

```
kubectl create secret tls myapp-tls --cert=tls.crt --key=tls.key
```

TLS Ingress:

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: tls-ingress
spec:
  tls:
  - hosts:
    - myapp.local
    secretName: myapp-tls
  rules:
  - host: myapp.local
    http:
      paths:
      - path: /
        pathType: Prefix
      backend:
        service:
          name: app1
          port:
            number: 80

```

---

## Exam Practice

### Scenario 1

Create Ingress for path /web routing to webapp service on port 80.

Solution

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: webapp-ingress
spec:
  rules:
  - http:
      paths:
      - path: /web
        pathType: Prefix
        backend:
          service:
            name: webapp
            port:
              number: 80
```

---

### □ Cleanup

```
kubectl delete ingress --all
kubectl delete deployment app1 app2
kubectl delete svc app1 app2
kubectl delete secret myapp-tls --ignore-not-found
rm -f tls.key tls.crt
```

---

## What We Learned

- ☒ Ingress resource creation
  - ☒ Path-based routing
  - ☒ Host-based routing
  - ☒ pathType (Prefix, Exact)
  - ☒ TLS configuration
- 

Lab 11 | Lab 13: Debugging # Lab 13: Debugging & Troubleshooting

## Learning Objectives

- Identify pod issues
  - Debugging commands
  - Common issues and solutions
-

## Debugging Flow




Figure 24: Mermaid Diagram

---

## Hands-on Exercises

### Exercise 1: Pod States

Status	Meaning	Solution
Pending	Can't schedule	Check resources, taints
ContainerCreating	Pulling image	Wait or check image
Running	Working	-
CrashLoopBackOff	Keeps crashing	Check logs
ImagePullBackOff	Can't pull image	Check image name
Error	Error occurred	Check describe/logs

---

### Exercise 2: kubectl describe

Solution

```
# Create a broken pod
kubectl run broken --image=nginx:wrong-tag
```

```
kubectl describe pod broken
# Look at Events section
```

---

### **Exercise 3: kubectl logs**

Solution

```
# Current logs  
kubectl logs <pod>  
  
# Previous (crashed) container  
kubectl logs <pod> --previous  
  
# Follow logs  
kubectl logs -f <pod>  
  
# Last N lines  
kubectl logs --tail=20 <pod>  
  
# Multi-container pod  
kubectl logs <pod> -c <container>
```

---

### **Exercise 4: kubectl exec**

Solution

```
# Run a command  
kubectl exec <pod> -- ls -la /  
  
# Interactive shell  
kubectl exec -it <pod> -- /bin/sh  
  
# Network debug  
kubectl exec <pod> -- curl -s localhost:80  
kubectl exec <pod> -- nslookup kubernetes
```

---

### **Exercise 5: Events**

Solution

```
# All events  
kubectl get events --sort-by='.lastTimestamp'  
  
# Warnings only  
kubectl get events --field-selector type=Warning  
  
# Specific pod  
kubectl get events --field-selector involvedObject.name=<pod>
```

---

### **Exercise 6: Network Debugging**

Solution

```

# Check service endpoints
kubectl get endpoints <service>

# DNS test
kubectl run dns-test --image=busybox --rm -it --restart=Never -- nslookup <service>

# Service access test
kubectl run test --image=busybox --rm -it --restart=Never -- wget -qO- <service>:<port>

```

---

## Common Issues Cheatsheet

### ImagePullBackOff

```

# Check
kubectl describe pod <pod> | grep -A3 Events

# Solutions:
# 1. Wrong image name
# 2. Private registry - imagePullSecrets missing
# 3. Tag doesn't exist

```

### CrashLoopBackOff

```

# Check
kubectl logs <pod> --previous

# Solutions:
# 1. Wrong command/args
# 2. App error
# 3. Liveness probe too aggressive

```

### Pending

```

# Check
kubectl describe pod <pod>

# Solutions:
# 1. Insufficient resources
# 2. Node selector/affinity
# 3. Taints/tolerations
# 4. PVC not bound

```

---

## Exam Practice

### Scenario 1

web-pod is running but webpage not loading. Debug it.

#### Solution

```

# 1. Pod status
kubectl get pod web-pod

```

```
# 2. Describe
kubectl describe pod web-pod

# 3. Logs
kubectl logs web-pod

# 4. Test from inside
kubectl exec web-pod -- curl localhost:80

# 5. Check service
kubectl get svc
kubectl get endpoints
```

---

## □ Cleanup

```
kubectl delete pod broken --ignore-not-found
```

---

## What We Learned

- ☒ Understanding pod states
  - ☒ kubectl describe
  - ☒ kubectl logs (-previous, -f)
  - ☒ kubectl exec
  - ☒ kubectl get events
  - ☒ Common issues and solutions
- 

Lab 12 | Lab 14: Helm # Lab 14: Helm Basics

## Learning Objectives

- Understand Helm
  - Install charts
  - Manage releases
  - Basic Helm commands
- 

## What is Helm?



Figure 25: Mermaid Diagram

Concept	Description
<b>Chart</b>	Kubernetes app package
<b>Release</b>	Chart installation
<b>Repository</b>	Chart storage
<b>Values</b>	Configuration

---

## Hands-on Exercises

### Exercise 1: Add Repository

Solution

```
helm repo add bitnami https://charts.bitnami.com/bitnami
helm repo list
helm repo update
```

---

### Exercise 2: Search Charts

Solution

```
helm search hub nginx
helm search repo nginx
helm search repo bitnami/nginx --versions
```

---

### Exercise 3: Chart Info

Solution

```
helm show chart bitnami/nginx
helm show values bitnami/nginx
helm show all bitnami/nginx
```

---

### Exercise 4: Install Chart

Solution

```
helm install my-nginx bitnami/nginx

# With namespace
helm install my-nginx bitnami/nginx -n web --create-namespace

# Dry-run
helm install my-nginx bitnami/nginx --dry-run
```

---

## Exercise 5: Custom Values

Solution

```
# CLI value
helm install my-nginx bitnami/nginx --set replicaCount=3

# Values file
cat <<EOF > my-values.yaml
replicaCount: 2
service:
  type: ClusterIP
  port: 8080
EOF

helm install my-nginx bitnami/nginx -f my-values.yaml
```

---

## Exercise 6: Release Management

Solution

```
helm list
helm list -A
helm status my-nginx
helm history my-nginx
```

---

## Exercise 7: Upgrade and Rollback

Solution

```
helm upgrade my-nginx bitnami/nginx --set replicaCount=5
helm history my-nginx
helm rollback my-nginx 1
```

---

## Exercise 8: Uninstall

Solution

```
helm uninstall my-nginx
helm uninstall my-nginx -n web
```

---

## Helm Commands Summary

Command	Description
helm repo add	Add repository
helm repo update	Update repo
helm search repo	Search charts
helm show values	Show config options

Command	Description
helm install	Install chart
helm upgrade	Upgrade release
helm rollback	Rollback
helm uninstall	Uninstall
helm list	List releases
helm template	Render YAML

---

## Exam Practice

### Scenario 1

Install redis chart from bitnami as my-cache.

Solution

```
helm repo add bitnami https://charts.bitnami.com/bitnami
helm install my-cache bitnami/redis
```

---

### Scenario 2

Upgrade my-cache to 3 replicas.

Solution

```
helm upgrade my-cache bitnami/redis --set replica.replicaCount=3
```

---

### □ Cleanup

```
helm uninstall my-nginx --ignore-not-found
helm uninstall my-cache --ignore-not-found
rm -f my-values.yaml
```

---

## What We Learned

- ☒ Helm repository management
  - ☒ Chart search and info
  - ☒ helm install/upgrade/rollback/uninstall
  - ☒ Custom values
  - ☒ Release management
- 

Lab 13 | Lab 15: Deployment Strategies # Lab 15: Blue-Green & Canary Deployments

## Learning Objectives

- Blue-Green deployment strategy
  - Canary deployment strategy
  - Traffic management with service selector
- 

## Deployment Strategies




Figure 26: Mermaid Diagram

Strategy	Description	Risk
<b>Blue-Green</b>	Instant switch, old version standby	Low
<b>Canary</b>	Gradual rollout, test with small traffic	Very low
<b>Rolling</b>	Default K8s, gradual update	Medium

---

## Blue-Green Deployment

### Exercise 1: Blue Deployment

Solution

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-blue
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
      version: blue
  template:
    metadata:
      labels:
        app: myapp
        version: blue
  spec:
```

```
  containers:
  - name: app
    image: nginx:1.19
```

---

## Exercise 2: Service (Point to Blue)

Solution

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-svc
spec:
  selector:
    app: myapp
    version: blue  # Points to blue
  ports:
  - port: 80
```

---

## Exercise 3: Green Deployment

Solution

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-green
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
      version: green
  template:
    metadata:
      labels:
        app: myapp
        version: green
    spec:
      containers:
      - name: app
        image: nginx:1.21
```

---

## Exercise 4: Switch Blue → Green

Solution

```
# Patch service selector
kubectl patch svc myapp-svc -p '{"spec":{"selector":{"version":"green"}}}'
```

Instant switch! All traffic goes to green.

---

## Canary Deployment

### Exercise 5: Canary Setup

10% traffic to canary.

Solution

```
# Stable: 9 replicas
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-stable
spec:
  replicas: 9
  selector:
    matchLabels:
      app: myapp
      track: stable
  template:
    metadata:
      labels:
        app: myapp
        track: stable
    spec:
      containers:
        - name: app
          image: nginx:1.19
---
# Canary: 1 replica
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-canary
spec:
  replicas: 1
  selector:
    matchLabels:
      app: myapp
      track: canary
  template:
    metadata:
      labels:
        app: myapp
        track: canary
    spec:
      containers:
        - name: app
          image: nginx:1.21
---
# Service selects both
apiVersion: v1
kind: Service
```

```
metadata:  
  name: myapp-svc  
spec:  
  selector:  
    app: myapp  # Selects both deployments  
  ports:  
  - port: 80
```

9 stable + 1 canary = 10% canary traffic

---

## Exercise 6: Scale Up Canary

Solution

```
kubectl scale deployment app-stable --replicas=5  
kubectl scale deployment app-canary --replicas=5  
# Now 50/50 traffic
```

---

## Exam Practice

### Scenario 1

web-blue deployment exists. Create web-green and switch web-svc to green.

Solution

```
kubectl create deployment web-green --image=nginx:1.21  
kubectl label deployment web-green version=green  
kubectl patch svc web-svc -p '{"spec": {"selector": {"version": "green"}}}'
```

---

### □ Cleanup

```
kubectl delete deployment app-blue app-green app-stable app-canary --ignore-not-found  
kubectl delete svc myapp-svc --ignore-not-found
```

---

## What We Learned

- Blue-Green deployment
  - Traffic management with service selector
  - Canary deployment
  - Instant vs gradual rollout
- 

Lab 14 | Lab 16: Kustomize # Lab 16: Kustomize

## Learning Objectives

- Understand Kustomize
- Base and overlay structure

- `kubectl kustomize`
  - Patches and transformations
- 

## What is Kustomize?




Figure 27: Mermaid Diagram

**Kustomize** customizes YAML without templates:  
- Built into `kubectl` (`kubectl apply -k`)  
- Alternative to Helm  
- Base + Overlay structure

---

## Hands-on Exercises

### Setup: Directory Structure

```
mkdir -p kustomize-demo/{base,overlays/dev,overlays/prod}  
cd kustomize-demo
```

---

### Exercise 1: Create Base

Solution

```
# base/deployment.yaml  
cat <<EOF > base/deployment.yaml  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: myapp  
spec:  
  replicas: 1  
  selector:  
    matchLabels:  
      app: myapp  
  template:  
    metadata:  
      labels:  
        app: myapp  
    spec:  
      containers:  
      - name: app  
        image: nginx:1.19  
EOF  
  
# base/kustomization.yaml  
cat <<EOF > base/kustomization.yaml  
apiVersion: kustomize.config.k8s.io/v1beta1  
kind: Kustomization  
resources:  
- deployment.yaml  
EOF
```

---

### Exercise 2: Dev Overlay

Solution

```
cat <<EOF > overlays/dev/kustomization.yaml  
apiVersion: kustomize.config.k8s.io/v1beta1  
kind: Kustomization  
  
resources:  
- ../../base  
  
namePrefix: dev-
```

```
replicas:
- name: myapp
  count: 1

commonLabels:
  env: development
EOF
```

---

### Exercise 3: Prod Overlay

Solution

```
cat <<EOF > overlays/prod/kustomization.yaml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

resources:
- .../base

namePrefix: prod-

replicas:
- name: myapp
  count: 5

commonLabels:
  env: production

images:
- name: nginx
  newTag: "1.21"
EOF
```

---

### Exercise 4: Build and Preview

Solution

```
# Preview dev
kubectl kustomize overlays/dev

# Preview prod
kubectl kustomize overlays/prod

# Save to file
kubectl kustomize overlays/prod > prod-manifests.yaml
```

---

### Exercise 5: Apply with Kustomize

Solution

```
# Apply with -k flag
kubectl apply -k overlays/dev

# Check
kubectl get all -l env=development
```

---

## Exercise 6: Patches

Solution

```
# overlays/prod/memory-patch.yaml
cat <<EOF > overlays/prod/memory-patch.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp
spec:
  template:
    spec:
      containers:
      - name: app
        resources:
          limits:
            memory: "256Mi"
EOF

# Add to kustomization.yaml
cat <<EOF >> overlays/prod/kustomization.yaml

patches:
- path: memory-patch.yaml
EOF
```

---

## Kustomization.yaml Options

Option	Description
resources	Base YAML files
namePrefix	Add prefix to all names
nameSuffix	Add suffix to all names
namespace	Set namespace
commonLabels	Add labels to all resources
images	Change image tags
replicas	Change replica count
patches	Strategic merge patches
configMapGenerator	Generate ConfigMaps

---

## Exam Practice

### Scenario 1

Apply base/ with namespace production.

Solution

```
cat <<EOF > overlay/kustomization.yaml
resources:
- ../base
namespace: production
EOF
```

```
kubectl apply -k overlay/
```

---

### □ Cleanup

```
kubectl delete -k overlays/dev --ignore-not-found
kubectl delete -k overlays/prod --ignore-not-found
cd .. && rm -rf kustomize-demo
```

---

## What We Learned

- Base and overlay structure
  - kustomization.yaml
  - kubectl kustomize / kubectl apply -k
  - Patches
  - ConfigMap generators
- 

Lab 15 | Lab 17: Dockerfile # Lab 17: Dockerfile & Container Basics

## Learning Objectives

- Write Dockerfiles
  - Build container images
  - Multi-stage builds
  - Security best practices
- 

## Container Structure

---

## Hands-on Exercises

### Exercise 1: Simple Dockerfile

Solution




Figure 28: Mermaid Diagram  
81

```

mkdir docker-demo && cd docker-demo

# index.html
cat <<EOF > index.html
<!DOCTYPE html>
<html>
<head><title>CKAD App</title></head>
<body><h1>Hello CKAD!</h1></body>
</html>
EOF

# Dockerfile
cat <<EOF > Dockerfile
FROM nginx:alpine
COPY index.html /usr/share/nginx/html/
EXPOSE 80
CMD ["nginx", "-g", "daemon off;"]
EOF

```

---

## Exercise 2: Dockerfile Directives

Directive	Description
FROM	Base image
WORKDIR	Working directory
COPY	Copy files
ADD	Copy files (tar extract, URL)
RUN	Build-time command
CMD	Default command
ENTRYPOINT	Fixed command
ENV	Environment variable
EXPOSE	Port documentation
USER	Run as user
ARG	Build argument

---

## Exercise 3: Python App Dockerfile

Solution

```

FROM python:3.9-slim

WORKDIR /app

COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

COPY app.py .

EXPOSE 5000

```

```
USER 1000  
CMD ["python", "app.py"]
```

---

#### Exercise 4: Multi-Stage Build



Figure 29: Mermaid Diagram

#### Solution

```
# Build stage  
FROM golang:1.19 AS builder  
WORKDIR /app  
COPY main.go .  
RUN CGO_ENABLED=0 go build -o myapp main.go  
  
# Runtime stage  
FROM alpine:3.18  
WORKDIR /app  
COPY --from=builder /app/myapp .  
USER 1000  
CMD ["/myapp"]
```

Multi-stage benefits: - Smaller final image - Build tools not in runtime - Better security

---

#### Exercise 5: Security Best Practices

##### Solution

```
# Good Dockerfile  
FROM python:3.9-slim  
  
# Non-root user  
RUN useradd -r -u 1000 appuser  
  
WORKDIR /app  
  
# Only necessary files  
COPY requirements.txt .  
RUN pip install --no-cache-dir -r requirements.txt  
  
COPY --chown=appuser:appuser app.py .
```

```
USER appuser
EXPOSE 5000
CMD ["python", "app.py"]
```

Best practices: - Use non-root user - Minimal base image (alpine, slim) - --no-cache with pip/apt  
- Specific version tags - Don't use latest tag - Don't run as root

---

## Exercise 6: .dockerignore

Solution

```
cat <<EOF > .dockerignore
.git
.gitignore
Dockerfile
*.md
__pycache__
*.pyc
.env
node_modules
EOF
```

---

## Exercise 7: K3s Usage (containerd)

K3s uses containerd. To import images:

Solution

```
# Build with Docker, export as tar
docker build -t myapp:v1 .
docker save myapp:v1 -o myapp.tar

# Import to K3s
sudo k3s ctr images import myapp.tar

# Use in pod
kubectl run myapp --image=myapp:v1 --image-pull-policy=Never
```

---

## Exam Practice

### Scenario 1

Fix this Dockerfile:

```
FROM ubuntu:latest
COPY . .
RUN apt-get install python3
CMD python3 app.py
```

Solution

```
FROM python:3.9-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
COPY app.py .
USER 1000
CMD ["python3", "app.py"]
```

Issues fixed: - ubuntu:latest → specific slim image - Missing apt-get update - Missing WORKDIR - Missing non-root user - CMD should use exec form

---

## □ Cleanup

```
cd .. && rm -rf docker-demo
```

---

## What We Learned

- ☒ Dockerfile directives
  - ☒ Multi-stage builds
  - ☒ Security best practices
  - ☒ K3s containerd integration
  - ☒ .dockerignore
- 

[Lab 16](#) | [Home](#)